

NARENDRANATH EDARA

+1 (573) 466-6656 | edara.narendranath@gmail.com | linkedin.com/in/narendranathe | github.com/narendranathe | narendranathe.github.io

Data Platforms & Applied AI. Owns payroll data systems and tools teams adopt. Cut refresh from 30 minutes to under 8 and freed 5,000 engineering hours annually; deployed schema-aware AI assistance for safer SQL.

TECHNICAL SKILLS

Data Engineering: Python, SQL, ETL, data modeling, data quality, CDC, T-SQL, SQL Server, Microsoft Fabric, Power BI
Applied AI & ML: Schema context, Claude Code CLI, GitHub Copilot, NLP, RAG, pgvector, PyTorch, QLoRA, FastAPI
Distributed Data: Apache Spark, Structured Streaming, Kafka, Delta Lake, Delta tables, lakehouse architecture
Cloud & DevOps: Azure, Azure DevOps, CI/CD, AKS, Docker, SQL Agent, SSISDB, Always On Availability Groups

EXPERIENCE

Data Engineer | *ExponentHR*

July 2024 - Present | Dallas, TX

Owns payroll ingestion, recovery and releases; builds team tooling and onboarding documentation used by new hires.

- Standardized SQL Server incremental loading with change data capture (CDC) and idempotent MERGE: cut refresh from 30 minutes to **under 8 minutes**, lowered compute cost 67%, and made failed loads safe to rerun.
- Freed **5,000 engineering hours** annually by automating 20 daily copy/restore runs on Contained Always On Availability Groups; made retries safe by resolving orphaned CDC jobs, replication races and stuck restores.
- Deployed an **AI context harness** the team uses for safer SQL and faster bug resolution: schema-derived primary/foreign keys and table grain, annotated with business meaning, guide Claude Code CLI during code changes.
- Cut support tickets **40%** and report queries from 12 seconds to under 4 through Microsoft Fabric semantic models, natural-language-to-SQL access, and SQL Server index and stored-procedure tuning on reporting workloads.
- Cut releases from 3 months to **14 days** with Azure DevOps CI/CD, automated reporting releases and code-review gates; Python delivery reporting over Git and Azure DevOps APIs saved 15+ engineering hours per sprint.
- Made SQL Agent failures diagnosable with an SSISDB procedure logging package names, run times and execution evidence; control-table validation exposes data drift so the team can investigate and ship targeted fixes.

Data Engineer | *Missouri S&T*

Aug. 2023 - July 2024 | Rolla, MO

University research and cloud infrastructure: built NLP sentiment analysis and Azure AI anomaly detection pipelines.

- Coauthored a **peer-reviewed NLP** study using Python, VADER and RoBERTa on TripAdvisor and Yelp museum reviews; compared sentiment with visitor ratings to identify how data sources and methods affect conclusions.
- Caught a production memory leak 4 hours before an outage with Azure AI Anomaly Detector pipelines that chose an algorithm per time-series profile, reaching more than 95% accuracy across research workloads.
- Cut Azure spend \$3,200 a month by moving anomaly detection from 20 static virtual machines to Azure Kubernetes Service with horizontal pod autoscaling, lifting CPU utilization from 12% to 64% on 4 to 8 nodes.

Product Intern | *C2FO*

June 2023 - Aug. 2023 | Kansas City, KS

B2B fintech working-capital transaction platform: summer product internship working with product and operations.

- Analyzed transactions and user behavior with SQL to guide product decisions on a working-capital platform.
- Halved resource-allocation time by writing preferred-offers product requirement documents from stakeholder interviews.

Business Intelligence Analyst | *Udaan*

Sept. 2020 - Mar. 2021 | Nashik, India

Nashik B2B staples launch: owned gross merchandise value, demand planning, procurement and sourcing-team delivery.

- Grew Nashik gross merchandise value to INR 50 million per month during the Staples launch, matching the incumbent; owned demand planning and procurement, reaching INR 39 million in December at 1.3x month-over-month growth.
- Hired and managed a 10+ person sourcing team; sustained 99.3% order fulfillment through demand-based capacity planning and kept delivery cycles running through application downtime with a business continuity plan.

Business Analyst | *Zomato*

Mar. 2018 - Sept. 2020 | *Hyderabad, India*

Food-delivery marketplace across South India: owned pricing intelligence, search analysis, and a 300-restaurant portfolio.

- Turned **contribution per order** from negative 18 to positive 2 rupees across 300 restaurants by rebalancing discount sharing against advertising revenue, moving the portfolio from loss to positive unit economics.
- Contributed to a 9% market-share gain in contested metros by building competitor analytics on price, delivery time, and coverage, with ingestion and refresh that kept pricing on current rather than week-old signals.
- Cut support-desk volume 80% by indexing 100,000+ internal documents in Elasticsearch for self-service search.

PROJECTS

repo-context-hooks | *Python, GitHub Actions, Sigstore, CodeQL*

PyPI

- Published a zero-dependency Python package that snapshots repository context around each commit, letting LLM coding agents resume work across sessions while keeping downstream projects free of added runtime dependencies.
- Secured publishing with Sigstore signatures and CodeQL checks; validated the release with **330+ passing tests**.

FinTune | *Mistral-7B, QLoRA, FastAPI, PyTorch*

2026

- Fine-tuned Mistral-7B with QLoRA (4-bit NF4) for financial sentiment and served it through FastAPI with PII redaction before inference and KL-divergence drift monitoring; 35+ tests cover data, model, and serving.
- Automated recovery with a three-state circuit breaker: model reload, smaller batches, or quantized fallback.

AutoApply AI | *Chrome MV3, FastAPI, pgvector, React*

2026

- Built a privately deployed Chrome extension and FastAPI service using pgvector retrieval-augmented generation (RAG) to ground answers in work history; supports six LLM providers and 11 ATS adapters, with **355 backend tests**.

Portfolio Risk Analytics (in progress) | *Kafka, Spark Structured Streaming, Delta Lake, FastAPI*

2026

- Built Kafka-to-Spark Structured Streaming ingestion of market ticks into a Delta Lake bronze layer with checkpointing, 5-second windows, and a 10-second watermark; the FastAPI value-at-risk handoff is still in progress.

EDUCATION & CERTIFICATIONS

Missouri S&T - MS Information Science and Technology, GPA 4.0/4.0

Jan. 2022 - Dec. 2023

GVP College of Engineering - B.Tech Mechanical Engineering

2017

DP-700 Microsoft Certified: Fabric Data Engineer Associate | Certified Scrum Product Owner (CSPO)

Publication: *Journal of Nonprofit & Public Sector Marketing*, 2025, DOI 10.1080/10495142.2025.2525123